

Practical 1: Data Wrangling

Nina van Gerwen

13-09-2018

First, load the packages:

```
library(ISLR)
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.2 --
## v ggplot2 3.3.6      v purrr   0.3.4
## v tibble  3.1.6      v dplyr   1.0.8
## v tidyr   1.2.0      v stringr 1.4.0
## v readr   2.1.2      v forcats 0.5.1
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()

library(haven)
library(readxl)
```

1. Objects and classes

```
object_1 <- 1:5 ## integers
object_2 <- 1L:5L ## integers
object_3 <- "-123.456" ## character
object_4 <- as.numeric(object_2) ## numeric
object_5 <- letters[object_1] ## characters
object_6 <- as.factor(rep(object_5, 2)) ## factors
object_7 <- c(1, 2, 3, "4", "5", "6") ## characters
```

2. Converting

```
object_7 <- as.numeric(object_7)
```

3. Making a list

```
objects <- list(object_1, object_2, object_3, object_4, object_5,
                object_6, object_7)
```

4. Making a data frame

```
object_frame <- data.frame(object_1, object_2, object_5)
```

5. Determining the size of a data frame

```
ncol(object_frame) ## 3
```

```
## [1] 3
```

```
nrow(object_frame) ## 5
```

```
## [1] 5
```

Loading, viewing and summarising data

6. Importing google data

```
apps <- read_csv("Data/googleplaystore.csv")
```

```
## Rows: 10841 Columns: 13
## -- Column specification -----
## Delimiter: ","
## chr (11): App, Category, Size, Installs, Type, Price, Content Rating, Genres...
## dbl (2): Rating, Reviews
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

7. Unexpected results

```
str(apps)
```

```
## spec_tbl_df [10,841 x 13] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## $ App          : chr [1:10841] "Photo Editor & Candy Camera & Grid & ScrapBook" "Co
## $ Category     : chr [1:10841] "ART_AND_DESIGN" "ART_AND_DESIGN" "ART_AND_DESIGN" "
## $ Rating       : num [1:10841] 4.1 3.9 4.7 4.5 4.3 4.4 3.8 4.1 4.4 4.7 ...
## $ Reviews      : num [1:10841] 159 967 87510 215644 967 ...
```

```
## $ Size      : chr [1:10841] "19M" "14M" "8.7M" "25M" ...
## $ Installs   : chr [1:10841] "10,000+" "500,000+" "5,000,000+" "50,000,000+" ...
## $ Type       : chr [1:10841] "Free" "Free" "Free" "Free" ...
## $ Price      : chr [1:10841] "0" "0" "0" "0" ...
## $ Content Rating: chr [1:10841] "Everyone" "Everyone" "Everyone" "Teen" ...
## $ Genres     : chr [1:10841] "Art & Design" "Art & Design;Pretend Play" "Art & De
## $ Last Updated : chr [1:10841] "January 7, 2018" "January 15, 2018" "August 1, 2018
## $ Current Ver  : chr [1:10841] "1.0.0" "2.0.0" "1.2.4" "Varies with device" ...
## $ Android Ver  : chr [1:10841] "4.0.3 and up" "4.0.3 and up" "4.0.3 and up" "4.2 an
## - attr(*, "spec")=
## .. cols(
## ..   App = col_character(),
## ..   Category = col_character(),
## ..   Rating = col_double(),
## ..   Reviews = col_double(),
## ..   Size = col_character(),
## ..   Installs = col_character(),
## ..   Type = col_character(),
## ..   Price = col_character(),
## ..   `Content Rating` = col_character(),
## ..   Genres = col_character(),
## ..   `Last Updated` = col_character(),
## ..   `Current Ver` = col_character(),
## ..   `Android Ver` = col_character()
## .. )
## - attr(*, "problems")=<externalptr>
```

The variables 'Category', 'Size', 'Installs', 'Price', 'Content Rating', 'Genres' and 'Android Ver' should probably be a Factor instead of a Character.

8. The head

```
head(apps)
```

```
## # A tibble: 6 x 13
##   App          Category Rating Reviews Size  Installs Type  Price `Content Rating`
##   <chr>         <chr>    <dbl>   <dbl> <chr> <chr>    <chr> <chr> <chr>
## 1 "Photo Ed~ ART_AND~    4.1     159 19M   10,000+ Free  0     Everyone
## 2 "Coloring~ ART_AND~    3.9     967 14M   500,000+ Free  0     Everyone
## 3 "U Launch~ ART_AND~    4.7    87510 8.7M   5,000,0~ Free  0     Everyone
## 4 "Sketch -- ART_AND~    4.5   215644 25M   50,000,~ Free  0     Teen
## 5 "Pixel Dr~ ART_AND~    4.3     967 2.8M   100,000+ Free  0     Everyone
## 6 "Paper fl~ ART_AND~    4.4     167 5.6M   50,000+ Free  0     Everyone
## # ... with 4 more variables: Genres <chr>, `Last Updated` <chr>,
## #   `Current Ver` <chr>, `Android Ver` <chr>
```

9. Students data

```
## Loading data
students <- read_xlsx("Data/students.xlsx")

## Seeing the class of every variable
str(students)

## tibble [37 x 3] (S3: tbl_df/tbl/data.frame)
## $ student_number: num [1:37] 5117250 6562582 6000241 4862862 6561723 ...
## $ grade          : num [1:37] 6.54 7.57 6.08 7.71 6.57 ...
## $ programme      : chr [1:37] "A" "A" "B" "A" ...

head(students)

## # A tibble: 6 x 3
##   student_number grade programme
##         <dbl> <dbl> <chr>
## 1      5117250  6.54 A
## 2      6562582  7.57 A
## 3      6000241  6.08 B
## 4      4862862  7.71 A
## 5      6561723  6.57 B
## 6      5625916  7.90 B
```

Again, the variable 'Programme' should probably be a factor.

10. Summarising

```
summary(students)

## student_number      grade      programme
## Min.   :4011659    Min.   :4.844  Length:37
## 1st Qu.:4862862    1st Qu.:6.390  Class :character
## Median :6000241    Median :7.151  Mode  :character
## Mean   :5686729     Mean   :6.991
## 3rd Qu.:6553913    3rd Qu.:7.573
## Max.   :6997130     Max.   :9.291
```

The grades achieved by students ranges from 4.84 to 9.29 with an average grade of 7.15.

Data transformation

11. Filtering

```
students %>% filter(., grade < 5.5)
```

```
## # A tibble: 3 x 3
##   student_number grade programme
##         <dbl> <dbl> <chr>
## 1      6114656  5.16 A
## 2      5265402  5.49 B
## 3      4639846  4.84 A
```

12. More filtering

```
students %>% filter(., grade > 8 & programme == "A")
```

```
## # A tibble: 5 x 3
##   student_number grade programme
##         <dbl> <dbl> <chr>
## 1      6352581  8.09 A
## 2      6165611  8.02 A
## 3      4133949  8.40 A
## 4      4011659  8.94 A
## 5      6553913  8.24 A
```

13. Arranging

```
students %>% arrange(., programme, desc(grade))
```

```
## # A tibble: 37 x 3
##   student_number grade programme
##         <dbl> <dbl> <chr>
## 1      4011659  8.94 A
## 2      4133949  8.40 A
## 3      6553913  8.24 A
## 4      6352581  8.09 A
## 5      6165611  8.02 A
## 6      6997130  7.75 A
## 7      4862862  7.71 A
## 8      6562582  7.57 A
## 9      4483974  7.46 A
## 10     5128923  7.26 A
## # ... with 27 more rows
```

We use `desc(grade)` because we want the order to be descending and not ascending.

14. Selecting

```
students %>% select(., student_number, programme)
```

```
## # A tibble: 37 x 2
##   student_number programme
##         <dbl> <chr>
## 1      5117250 A
## 2      6562582 A
## 3      6000241 B
## 4      4862862 A
## 5      6561723 B
## 6      5625916 B
## 7      4096023 A
## 8      6114656 A
## 9      5265402 B
## 10     5977188 B
## # ... with 27 more rows
```

15. Mutating

```
students_recoded <- students %>% mutate(., programme = recode(programme,
                                                                A = "Science", B = "Social Science"))
```

So the above code states: students_recoded will be students, where we mutated the column programme such that A = Science, B = Social Science.

Data processing pipelines

16. Pipelining

```
popular_apps <-
  read_csv("Data/googleplaystore.csv") %>%
  mutate(Downloads = parse_number(Installs)) %>%
  filter(Downloads > 500000000) %>%
  arrange(desc(Rating)) %>%
  select(Rating, Category, Price, App) %>%
  distinct(App, .keep_all = TRUE)
```

```
## Rows: 10841 Columns: 13
## -- Column specification -----
## Delimiter: ","
## chr (11): App, Category, Size, Installs, Type, Price, Content Rating, Genres...
## dbl (2): Rating, Reviews
```

```
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
popular_apps
```

```
## # A tibble: 20 x 4
##   Rating Category      Price App
##   <dbl> <chr>      <chr> <chr>
## 1   4.5 GAME          0 "Subway Surfers"
## 2   4.5 SOCIAL        0 "Instagram"
## 3   4.5 PHOTOGRAPHY    0 "Google Photos"
## 4   4.4 COMMUNICATION  0 "WhatsApp Messenger"
## 5   4.4 TOOLS          0 "Google"
## 6   4.4 PRODUCTIVITY   0 "Google Drive"
## 7   4.3 COMMUNICATION  0 "Google Chrome: Fast & Secure"
## 8   4.3 COMMUNICATION  0 "Gmail"
## 9   4.3 ENTERTAINMENT   0 "Google Play Games"
## 10  4.3 TRAVEL_AND_LOCAL 0 "Maps - Navigate & Explore"
## 11  4.3 VIDEO_PLAYERS    0 "YouTube"
## 12  4.2 SOCIAL          0 "Google+"
## 13  4.2 TRAVEL_AND_LOCAL 0 "Google Street View"
## 14  4.1 COMMUNICATION    0 "Skype - free IM & video calls"
## 15  4.1 SOCIAL          0 "Facebook"
## 16  4   COMMUNICATION    0 "Messenger \x96 Text and Video Chat for Fre~
## 17  4   COMMUNICATION    0 "Hangouts"
## 18  3.9 BOOKS_AND_REFERENCE 0 "Google Play Books"
## 19  3.9 NEWS_AND_MAGAZINES 0 "Google News"
## 20  3.7 VIDEO_PLAYERS    0 "Google Play Movies & TV"
```

Grouping and summarisation

17 & 18. Summarising

The median of which variable exactly? I'm guessing Rating.

```
mad <- function(x){
  median(abs(x - median(x)))
}
popular_apps %>%
  summarise(
    median = median(Rating),
    min = min(Rating),
    max = max(Rating),
    mad = mad(Rating)
  )
```

```
## # A tibble: 1 x 4
##   median   min   max   mad
##   <dbl> <dbl> <dbl> <dbl>
## 1     4.3    3.7    4.5 0.150
```

19. Grouped summary

```
popular_apps %>%
  group_by(Category) %>%
  summarise(
    median = median(Rating),
    min = min(Rating),
    max = max(Rating),
    mad = mad(Rating)
  )
```

```
## # A tibble: 11 x 5
##   Category          median   min   max   mad
##   <chr>          <dbl> <dbl> <dbl> <dbl>
## 1 BOOKS_AND_REFERENCE 3.9    3.9    3.9 0
## 2 COMMUNICATION       4.2    4      4.4 0.150
## 3 ENTERTAINMENT       4.3    4.3    4.3 0
## 4 GAME                4.5    4.5    4.5 0
## 5 NEWS_AND_MAGAZINES  3.9    3.9    3.9 0
## 6 PHOTOGRAPHY         4.5    4.5    4.5 0
## 7 PRODUCTIVITY        4.4    4.4    4.4 0
## 8 SOCIAL              4.2    4.1    4.5 0.100
## 9 TOOLS               4.4    4.4    4.4 0
## 10 TRAVEL_AND_LOCAL   4.25   4.2    4.3 0.0500
## 11 VIDEO_PLAYERS      4      3.7    4.3 0.300
```

Final Exercise

Summary of how many Downloads per category:

```
own_summary <-
  read_csv("Data/googleplaystore.csv") %>%
  mutate(Downloads = parse_number(Installs)) %>%
  group_by(Category) %>%
  summarise(
    mean = mean(Downloads),
    min = min(Downloads),
    max = max(Downloads),
    var = var(Downloads)
```



```
) %>%
  arrange(desc(mean))
```

```
## Rows: 10841 Columns: 13
## -- Column specification -----
## Delimiter: ","
## chr (11): App, Category, Size, Installs, Type, Price, Content Rating, Genres...
## dbl (2): Rating, Reviews
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
own_summary
```

```
## # A tibble: 34 x 5
##   Category          mean    min      max    var
##   <chr>          <dbl> <dbl>    <dbl> <dbl>
## 1 COMMUNICATION 84359887.     1 1000000000 5.42e16
## 2 SOCIAL        47694467.     0 1000000000 3.07e16
## 3 VIDEO_PLAYERS 35554301.    10 1000000000 2.09e16
## 4 PRODUCTIVITY  33434178.     0 1000000000 1.52e16
## 5 GAME          30669602.     1 1000000000 9.40e15
## 6 PHOTOGRAPHY   30114172.     5 1000000000 1.24e16
## 7 TRAVEL_AND_LOCAL 26623594.     0 1000000000 1.92e16
## 8 NEWS_AND_MAGAZINES 26488755.     0 1000000000 1.88e16
## 9 ENTERTAINMENT  19256107. 10000 1000000000 7.18e15
## 10 TOOLS         13585732.     1 1000000000 5.25e15
## # ... with 24 more rows
```

From this summary, we can see that on average 'Communication' apps get the most downloads. However, we do not know whether the difference between the other apps categories and Communication is significant (though I would bet good money on it).